

# Signal Architecture Dashboard — Update Brief for Sig

February 18, 2026

---

## Versioning (Keep It Simple)

**Origin** — 1 agent, 1 pass, Signal was born. The generalist. **V1** — In production today. Geo-proximity framework, grouped ZIPs, multiple agents per group. Investors use this. **V2** — What we're building now. 10 domain specialists per ZIP, Layer 0 strategy education, Layer 2 intelligence enrichment, bias protection chain.

No other version numbers. No v3, v4, iteration labels anywhere on the site.

---

## Page-by-Page Updates

### 1. Homepage (Map to Mailbox)

**Phase 2 description — update:**

Per ZIP: Layer 0 strategy education → 10 parallel research agents → Layer 2 intelligence enrichment. Outputs include propensity to sell, mosaic patterns, conflict flags, and data quality assessment. All data SAVED.

**Phase 6 description — update:**

Layer 3 (Claude): First layer to see both research AND scoring columns. Maps intelligence to JS scoring functions. Layer 4 (Vercel): Executes against every row.

**"Key Insight" section** — Consider making "Research and scoring are decoupled" a first-class visual element, not a footnote. It's the single most important concept in the product.

---

### 2. Evolution Page

**Bottom tagline — change:**

- OLD: "The filtering principle at every layer of the architecture."

- NEW: Remove this line entirely. Just keep "All Signal. No Noise." The word "filtering" contradicts the enrichment architecture.

**V2 domain list — fix:** The 10th domain is listed as "narrative synthesis." It should be "micro-location intelligence." The actual 10 domains:

1. Physical Environment & Housing Stock
2. Financial Positions & Transaction Patterns
3. Distress & Legal Activity
4. Ownership Structure & Occupancy
5. Market Dynamics & Demographic Forces
6. Temporal & Seasonal Patterns
7. Economic Stress & Leading Indicators
8. Investor & Competitive Landscape
9. Regulatory & Political Environment
10. Neighborhood & Micro-Location Intelligence

**V2 agent count clarification:**

- Current: "500 Agents"
- Suggestion: Keep 500 but add context: "10 specialists × 50 ZIPs" so it doesn't read like 500 distinct agents

**V2 depth stat:**

- Current: "~30pg Per ZIP"
- Update to: "10–30pg Per ZIP" (varies by market data availability)

**Model attributions:**

- V2 uses Claude for Layer 0 and Layer 2, but Perplexity for the 10 research agents
- Consider noting this or just saying "Powered by Claude + Perplexity" for V2

**V1 agent count:**

- Current: "~8 Agents"
- More accurate: "3–10 Agents" (varied by geo group size)

### **3. Agents Page**

**Add a line above or below the agent grid explaining the philosophy:**

"Each agent receives strategy education and determines independently what matters within their domain. No prescribed fields. No form-filling. The agent decides what to research, how deep to go, and what's positive or negative for the investor's strategy."

**Agent descriptions** — current descriptions list expected data types like a schema (e.g., "Age, condition, vacancy"). These are fine as shorthand but slightly misleading since agents find what exists in each market, which may differ from what's listed. Optional: add a small note like "Representative focus areas — actual findings vary by market."

---

### **4. Bias Protection Page — Best Page, One Fix**

**Layer 2 card needs updating:**

Title change:

- OLD: "Intelligence Analyst" / "Strategy-blind synthesis & dedup"
- NEW: "Intelligence Enrichment" / "Strategy-blind enrichment & pattern recognition"

KNOWS section — replace:

- OLD: "Dedup logic"
- NEW: List should include: enrichment tagging, cross-agent confirmation, conflict detection, mosaic pattern recognition, proxy validation, propensity to sell assessment, data quality scoring

The key concept: Layer 2 doesn't condense or filter agent findings. Raw findings pass through the pipeline intact. Layer 2 ADDS intelligence on top — tags, patterns, validations, propensity. Information is gained, never lost.

Everything else on this page is accurate. The KNOWS / UNKNOWN / BLOCKED matrix is excellent.

---

### **5. Data Flow (Pipeline) Page**

**L2 node output description — update:**

- OLD: "deduped intel, conflicts, mosaic, propensity"
- NEW: "enrichment tags, conflicts, mosaic patterns, proxy validations, propensity, data

quality"

### Parse & Format Output node — update description:

- OLD: "IN: L2 analysis JSON / OUT: final intelligence package"
- NEW: "IN: raw agent findings (from Merge) + L2 enrichments / OUT: combined intelligence package"

This matters because the Parse node is the assembly point — it combines the raw agent findings that passed through the code pipeline with Layer 2's enrichment additions.

---

## 6. Status Page

### Move to "Built" (these are done):

- V2 workflow deploy & test ✓
- Full pipeline testing ✓ (all 10 agents validated, QA report generated)

### Add to "Planned":

- No-strategy mode (propensity-only for multi-strategy clients)
  - Contrasting market test (healthy/growing market validation)
  - Rerun consistency testing
  - Agent prompt tuning (Agents 5, 8, 10)
  - Multi-ZIP batch orchestration
- 

## Global Find-and-Replace

Across all pages, watch for and update these terms:

OLD	NEW
filtering / dedup / synthesis (when describing Layer 2)	enrichment / pattern recognition / validation
v3, v4, iteration numbers within V2	Remove — just "V2 architecture"
"Strategy-blind synthesis"	"Strategy-blind enrichment"

"deduped intel"

"enrichment tags"

## Summary

The dashboard structure is strong. These are content updates, not redesigns. The biggest theme is replacing "filtering/dedup/synthesis" language with "enrichment" language throughout — Layer 2 adds intelligence on top of agent findings, it never removes or condenses them. And keep versioning clean: Origin → V1 (production) → V2 (building now).